Project

Build some recommender systems.

Deliverables & Guidelines

The main deliverable is a GitHub repository with the following:

- A README file outlining the repository contents
- A requirements file with all software/package requirements to run your code
- A top level directory for Part I
    - Include a notebook or markdown with your approach and basic results
- A top level directory for Part II
    - Include a notebook or markdown with your approach and basic results

Projects should be completed in Python or R.

Data Scientists are more than statisticians and more than engineers. We solve real-life business problems by leveraging data and are usually brought in to brainstorm and frame a problem from day one. You will probably be expected to interface directly with other departments in the business, where your trust will be gained by sharing your insights into the core business problem at hand, and by translating these problems into a technical solution. This also means that you will be required to *communicate* your solution to key stake holders that are committed to solving the problem, but may not have your technical background.

Imagine a future employer looking at this repository. They would evaluate your project based on the solution's accuracy and your technical prowess, but they would also look for coherence and creativity. An employer would look for thoughtfulness and thoroughness; for example, how does the solution scale, were the hyper parameters tuned, will this solution discover novel recommendations, etc. Are the major contributions of your work clear? Did you call out important caveats?

These are the same standards on which you will be graded for these projects in this class.

Group instructions

You are organized into groups. Each team member will bring their own strengths and have their own opportunity areas for development. It's ok for teams to divide and conquer for divisible tasks, but *every team member should have a complete understanding of the entire solution, end to end*. I also recommend that before any piece solution is implemented by an individual, the entire team should brainstorm possible approaches, and should agree on the framework, what to try, and what the deliverable will look like.

Datasets

Datasets:
- https://gist.github.com/entaroadun/1653794
- https://gab41.lab41.org/the-nine-must-have-datasets-for-investigating-recommender-systems-ce9421bf981c
- http://www.recsyswiki.com/wiki/Category:Datasezt
- iHR - stay tuned

Part I - fundamentals

Due November 7th

Build two very simple collaborative filtering models. You may use published packages or methods - the goal of this exercise is to gain a practical intuition for how these types of common models work, and to develop methods to test and explore them.

- Choose a dataset / recommendation domain
- Treat this as a case study and think about it from a business perspective. What is your objective? What are you trying to optimizing and what are you willing to sacrifice?
- For this section, develop with a small dataset (< 10000 users / < 100 items; be thoughtful about how you sample your data)
- Build two brute-force collaborative filtering algorithms:
    1. Neighborhood-based (item *or* user)
    2. Model-based
- Develop evaluation methods:
    1. Cross-validation setup
    2. Accuracy on training and test data   Incorporate precision, recall
    3. Coverage on training and test data
- Systematically try a range of hyper parameters for your models (e.g. neighborhood size or number of latent dimensions). Record (*and explain* in your markdown) how your evaluation metrics change as a function of these parameters. *Include plots!*
- After seeing these results, what other design choices might you consider?
- How do the evaluation metrics change as a function of your model size? Systematically sample your data from a small size to a large size
    1. Does overall accuracy change?
    2. How does run-time scale with data size?   mention
- Referencing your case study set up from above, how do these methods meet your hypothetical objectives? Would you feel comfortable putting these solutions into production at a real company? What would be the potential watch outs?   add to conclusion of write-up

Evaluation

[10%] Case study framework   *Explain better in write-up*
  · How well was this envisioned as an actual business problem? Is it clear what is being solved, and what acceptance criteria would be?

[30%] Technical correctness   *Add a notebook 00, and requirements.txt*
  · Is the code complete and does it run?
  · Were appropriate methods applied in the right place? Were hyper-parameters used in a reasonable manner?

[10%] Evaluation methods   *coverage, precision, and recall*
  · Correct implementation

[20%] Model exploration
  · How well was each model explored?
    · Hyper-parameters
    · Data sample size
    · (Optional) Anything else?
  · Are plots clearly legible, and do they make sense?

[30%] Write up
  · If I were your manager at work, or a key stake holder, and I needed to understand what you did and why you did it in order to make a business decision (e.g. "go / no-go"), is this write up enough for me to go on? Would I have any lingering questions that were not already pointed out as "potential watch outs"?
  · Clarity matters
  · Connection to the business problem matters
  · Demonstrating an understanding of why the algorithms are performing or not, or *when* then do or do not, matters

---

Part II - deep dive

Due December 12th

Think of this part of the project as adding a bullet point or two to your resume, or adding a section to your public project portfolio. The outcome of this project is intended to be hosted in a public code repository where others can see your work, and most importantly, see how you think about a business problem, design a solution, and communicate your results.

1. Choose a personal/group objective:

(A) Gain practical experience with modern personalization frameworks and algorithms

Build a recommender system that goes beyond the standard CF techniques that you built from Part I. It is preferred that you build these algorithms mostly from scratch. Using previously published solutions (e.g. packages or blogs) as examples is ok, but try to avoid using black-box solutions (e.g. libFM), *except* if you are using them as a comparison to your solution, you are using them in a hybrid model, or you are building upon them in a novel way.

Some examples:

- Approximate nearest neighbors
- Content or hybrid-based
- Neural network based (wide & deep, RBM, VAE, etc)
- Recommender for novelty or serendipity (use iHR data)

Who should choose this option?

If you have limited previous experience with machine learning, are new to scientific computing, or are new to Python or R, this option is recommended. Is is also recommended if you'd like to explore new types of personalization algorithms, tie in knowledge that you've learned from other ML domains or classes to personalization, or try to build a kaggle-competitive solution.

(B) Gain practical experience with today's primary framework for distributed ML: Spark

Build a recommendation system using Spark. You can use PySpark, which can interact with a Spark context from within Python shell or instance, or Scala (less recommended - steep learning curve!).

Some examples:

- ALS collaborative filtering
  - Try different parameters, like regularization, non-negativity, or implicit vs. explicit
  - Grid search to find the best settings for your data
- Frequent pattern mining / Association rules
  - note: FPM is available for Python and for Data Frames, but Association rules is only available for Scala and Java, and only for RDD data types
- Locality Sensitive Hashing
  - Work on a large, sparse data set
    - Optimize the hyper-parameters

For all options above, how do the settings vary by data set? What about for popular items vs unpopular items, or power users versus infrequent users? How do they scale as you sample the data set?

You should benchmark all of these types of solutions against more traditional approaches, including a baseline, in terms of accuracy and time. You should also get creative. Consider engineering a new feature set to augment the features your data set came by using the data itself, or by matching the data to third-party data sets (e.g. "Does Wikipedia Information Help Netflix Predictions?").

Who should choose this option?

If you already have hands-on experience with algorithms beyond collaborative filtering this may be a valuable skill set to hone. Practical experience with Spark is currently a major selling point to hiring managers.

However, this option does carry some risk. Spark has a learning curve, and even getting simple things to run the first time can be tricky. **This is not a Spark class**; we will not be able to provide deep hands on help with Spark itself.

2. Choose a dataset

This could be the same or different than what you used in part I. Make sure that the data will support your project. You will need to do some minimum amount of ETL and exploration to make this determination.

3. Report

Along with the code, document your project with a notebook or markdown.

1. Clearly outline your objectives
2. Provide benchmarks. These should include a baseline bias model at the very least
3. Explore your model. Does it work equally well for all users and items? What about for less popular items or less prolific users? Think carefully about a business or technical framework to segment your data for users, items, or more, and test accuracy separately for these segments.
4. Devise methods to test for quality beyond just accuracy metrics. Consider testing coverage, novelty, serendipity etc
5. Extend your model. Can you think of any way to make your model better, like changing the objective, or adding in side information? It's ok if you can't make it better, but document your efforts and try to explain the results.

4. Evaluation

[30%] Technical correctness (see above)
[30%] Creativity
   · Choosing your model
   · Optimizing your model
   · Exploring your model
[40%] Presentation (see *writeup* above)
   · This is from the perspective of a potential employer. Most interviewers instinctively try to evaluate if the prospective employee will help them solve an outstanding or upcoming business problem, and will look for evidence that the candidate has already solved similar problems before. Does this project demonstrate that the author(s) would be able to frame and solve a sophisticated personalization problem at the hiring manager's company?